

1 GAN

```
import os, sys

# Add DLStudio-2.5.2 and AdversarialLearning to sys.path so Python can
# find DLStudio and AdversarialLearning
current_dir = os.getcwd()
print("current_dir = %s" % current_dir)

DLStudio_dir = os.path.join(current_dir, "../DLStudio-2.5.2")

sys.path.append(DLStudio_dir)

from DLStudio import *

current_dir = d:\MS Purdue\1.5\ECE60146\HW9
```

1.1 Model

I tried training my own model, but it collapsed lol. Additionally, I also tried default setting for CG2 in DLStudio; however, mode collapse occurred. Every generated image is an identical, abstract pattern that does not resemble human faces. Interestingly, the default setting for DG2 in DLStudio don't lead to mode collapse. However, when I **lowered** the learning rate for DG2, mode collapse started to occur — further investigation is needed.

```
# the code below is mainly borrowed from DLSudio-2.5.2

from DLStudio import DLStudio

import sys,os,os.path
import torch
import torch.nn as nn
import torch.nn.functional as F
import torchvision
import torchvision.transforms as tvf
import torchvision.transforms.functional as tvfF
import torch.optim as optim
import numpy as np
import math
import random
import matplotlib.pyplot as plt
import matplotlib.animation as animation
import time
import glob

import imageio
```

```

class MyAdversarialLearning(object):
    def __init__(self, *args, **kwargs ):
        if args:
            raise ValueError(
                '''AdversarialLearning constructor can only be
called with keyword arguments for the following
                keywords: epochs, learning_rate, batch_size,
momentum, image_size, dataroot, path_saved_model,
                use_gpu, latent_vector_size, ngpu, dlstudio,
device, LAMBDA, clipping_threshold, and beta1''')
            allowed_keys =
' dataroot', 'image_size', 'path_saved_model', 'momentum', 'learning_rate',
'epochs', 'batch_size', \
'classes', 'use_gpu', 'latent_vector_size', 'ngpu', 'dlstudio', 'beta1',
'LAMBDA', 'clipping_threshold'
            keywords_used = kwargs.keys()

            for keyword in keywords_used:

                if keyword not in allowed_keys:

                    raise SyntaxError(keyword + ": Wrong keyword used ---
check spelling")
                    learning_rate = epochs = batch_size = convo_layers_config =
momentum = None
                    image_size = fc_layers_config = dataroot = path_saved_model =
classes = use_gpu = None
                    latent_vector_size = ngpu = beta1 = LAMBDA =
clipping_threshold = None
                    if 'latent_vector_size' in kwargs :
latent_vector_size = kwargs.pop('latent_vector_size')
                    if 'ngpu' in kwargs : ngpu =
kwargs.pop('ngpu')
                    if 'dlstudio' in kwargs : dlstudio =
kwargs.pop('dlstudio')
                    if 'beta1' in kwargs : beta1 =
kwargs.pop('beta1')
                    if 'LAMBDA' in kwargs : LAMBDA =
kwargs.pop('LAMBDA')
                    if 'clipping_threshold' in kwargs :
clipping_threshold = kwargs.pop('clipping_threshold')
                    if latent_vector_size:
                        self.latent_vector_size = latent_vector_size
                    if ngpu:
                        self.ngpu = ngpu
                    if dlstudio:
                        self.dlstudio = dlstudio
                    if beta1:

```

```

        self.betal = betal
    if LAMBDA:
        self.LAMBDA = LAMBDA
    if clipping_threshold:
        self.clipping_threshold = clipping_threshold

    def show_sample_images_from_dataset(self, dlstudio):
        data = next(iter(self.train_dataloader))
        real_batch = data[0]

    self.dlstudio.display_tensor_as_image(torchvision.utils.make_grid(real
_batch, padding=2, pad_value=1, normalize=True))

    def set_dataloader(self):
        dataset =
torchvision.datasets.ImageFolder(root=self.dlstudio.dataroot,
                                transform = tvl.Compose([

tvt.Resize(self.dlstudio.image_size),
tvt.CenterCrop(self.dlstudio.image_size),
                                tvt.ToTensor(),
                                tvt.Normalize((0.5, 0.5,
0.5), (0.5, 0.5, 0.5)),
                                ]))
        self.train_dataloader = torch.utils.data.DataLoader(dataset,
batch_size=self.dlstudio.batch_size,
shuffle=True, num_workers=2)

    def weights_init(self,m):
        """
        Uses the DCGAN initializations for the weights
        """
        classname = m.__class__.__name__
        if classname.find('Conv') != -1:
            nn.init.normal_(m.weight.data, 0.0, 0.02)
        elif classname.find('BatchNorm') != -1:
            nn.init.normal_(m.weight.data, 1.0, 0.02)
            nn.init.constant_(m.bias.data, 0)

    def calc_gradient_penalty(self, netC, real_data, fake_data):
        """
        Implementation by Marvin Cao: https://github.com/caogang/wgan-gp

```

Marvin Cao's code is a PyTorch version of the Tensorflow based implementation provided by the authors of the paper "Improved Training of Wasserstein GANs" by Gulrajani, Ahmed, Arjovsky, Dumouli, and Courville.

```
"""
    BATCH_SIZE = self.dlstudio.batch_size
    LAMBDA = self.LAMBDA
    epsilon = torch.rand(1).cuda()
    interpolates = epsilon * real_data + ((1 - epsilon) *
fake_data)
    interpolates = interpolates.requires_grad_(True).cuda()
    critic_interpolates = netC(interpolates)
    gradients = torch.autograd.grad(outputs=critic_interpolates,
inputs=interpolates,

grad_outputs=torch.ones(critic_interpolates.size()).cuda(),
                                create_graph=True,
retain_graph=True, only_inputs=True)[0]
    gradient_penalty = ((gradients.norm(2, dim=1) - 1) **
2).mean() * LAMBDA
    return gradient_penalty
```

```
def close_event(self):
    """
    from stackoverflow.com
    """
    plt.close()
```

```
class DiscriminatorDG2(nn.Module):
    """
    This is essentially the same network as the DCGAN for DG1,
    except for the extra layer
    "self.extra" shown below. We also declare a batchnorm for
    this extra layer in the form
    of "self.bnX". In the implementation of "forward()", we
    invoke the extra layer at the
    beginning of the network.

    Class Path: AdversarialLearning -> DiscriminatorDG2
    """
    def __init__(self, skip_connections=True, depth=16):
        super(MyAdversarialLearning.DiscriminatorDG2,
self).__init__()
        self.conv_in = nn.Conv2d( 3, 64, kernel_size=4,
stride=2, padding=1)
        self.extra = nn.Conv2d( 64, 64, kernel_size=4,
stride=1, padding=2)
        self.conv_in2 = nn.Conv2d( 64, 128, kernel_size=4,
```

```

stride=2, padding=1)
self.conv_in3 = nn.Conv2d( 128, 256, kernel_size=4,
stride=2, padding=1)
self.conv_in4 = nn.Conv2d( 256, 512, kernel_size=4,
stride=2, padding=1)
self.conv_in5 = nn.Conv2d( 512, 1, kernel_size=4,
stride=1, padding=0)
self.bn1 = nn.BatchNorm2d(128)
self.bn2 = nn.BatchNorm2d(256)
self.bn3 = nn.BatchNorm2d(512)
self.bnX = nn.BatchNorm2d(64)
self.sig = nn.Sigmoid()
def forward(self, x):
    x = torch.nn.functional.leaky_relu(self.conv_in(x),
negative_slope=0.2, inplace=True)
    x = self.bnX(self.extra(x))
    x = torch.nn.functional.leaky_relu(x, negative_slope=0.2,
inplace=True)
    x = self.bn1(self.conv_in2(x))
    x = torch.nn.functional.leaky_relu(x, negative_slope=0.2,
inplace=True)
    x = self.bn2(self.conv_in3(x))
    x = torch.nn.functional.leaky_relu(x, negative_slope=0.2,
inplace=True)
    x = self.bn3(self.conv_in4(x))
    x = torch.nn.functional.leaky_relu(x, negative_slope=0.2,
inplace=True)
    x = self.conv_in5(x)
    x = self.sig(x)
    return x

class GeneratorDG2(nn.Module):
    """
    The Generator for DG2 is exactly the same as for the DG1. So
    please the comment block for that
    Generator.

    Class Path: AdversarialLearning -> GeneratorDG2
    """
    def __init__(self):
        super(MyAdversarialLearning.GeneratorDG2, self).__init__()
        self.latent_to_image = nn.ConvTranspose2d( 100, 512,
kernel_size=4, stride=1, padding=0, bias=False)
        self.upsampler2 = nn.ConvTranspose2d( 512, 256,
kernel_size=4, stride=2, padding=1, bias=False)
        self.upsampler3 = nn.ConvTranspose2d( 256, 128,
kernel_size=4, stride=2, padding=1, bias=False)
        self.upsampler4 = nn.ConvTranspose2d( 128, 64,
kernel_size=4, stride=2, padding=1, bias=False)

```

```

        self.upsampler5 = nn.ConvTranspose2d( 64, 3,
kernel_size=4, stride=2, padding=1, bias=False)
        self.bn1 = nn.BatchNorm2d(512)
        self.bn2 = nn.BatchNorm2d(256)
        self.bn3 = nn.BatchNorm2d(128)
        self.bn4 = nn.BatchNorm2d(64)
        self.tanh = nn.Tanh()
    def forward(self, x):
        x = self.latent_to_image(x)
        x = torch.nn.functional.relu(self.bn1(x))
        x = self.upsampler2(x)
        x = torch.nn.functional.relu(self.bn2(x))
        x = self.upsampler3(x)
        x = torch.nn.functional.relu(self.bn3(x))
        x = self.upsampler4(x)
        x = torch.nn.functional.relu(self.bn4(x))
        x = self.upsampler5(x)
        x = self.tanh(x)
        return x

##### DG2 Definition ENDS
#####

#####
#####
    ## The training routines follow, first for a GAN constructed
    using either the DG1 and or the DG2
    ## Discriminator-Generator Networks, and then for a WGAN
    constructed using either the CG1 or the CG2
    ## Critic-Generator Networks.

#####
#####

    def run_gan_code(self, dlstudio, discriminator, generator,
results_dir):
        """
        This function is meant for training a Discriminator-Generator
        based Adversarial Network.
        The implementation shown uses several programming constructs
        from the "official" DCGAN
        implementations at the PyTorch website and at GitHub.

        Regarding how to set the parameters of this method, see the
        following script

            dcgan_DG1.py

```

```

        in the "ExamplesAdversarialLearning" directory of the
distribution.
"""
    dir_name_for_results = results_dir
    if os.path.exists(dir_name_for_results):
        files = glob.glob(dir_name_for_results + "/*")
        for file in files:
            if os.path.isfile(file):
                os.remove(file)
            else:
                files = glob.glob(file + "/*")
                list(map(lambda x: os.remove(x), files))
    else:
        os.mkdir(dir_name_for_results)
    # Set the number of channels for the 1x1 input noise vectors
for the Generator:
    nz = 100
    netD = discriminator.to(self.dlstudio.device)
    netG = generator.to(self.dlstudio.device)
    # Initialize the parameters of the Discriminator and the
Generator networks according to the
    # definition of the "weights_init()" method:
    netD.apply(self.weights_init)
    netG.apply(self.weights_init)
    # We will use a the same noise batch to periodically check on
the progress made for the Generator:
    fixed_noise = torch.randn(self.dlstudio.batch_size, nz, 1, 1,
device=self.dlstudio.device)
    # Establish convention for real and fake labels during
training
    real_label = 1
    fake_label = 0
    # Adam optimizers for the Discriminator and the Generator:
    optimizerD = optim.Adam(netD.parameters(),
lr=dlstudio.learning_rate, betas=(self.betal, 0.999))
    optimizerG = optim.Adam(netG.parameters(),
lr=dlstudio.learning_rate, betas=(self.betal, 0.999))
    # Establish the criterion for measuring the loss at the
output of the Discriminator network:
    criterion = nn.BCELoss()
    # We will use these lists to store the results accumulated
during training:
    img_list = []
    G_losses = []
    D_losses = []
    iters = 0
    print("\n\nStarting Training Loop...\n\n")
    start_time = time.perf_counter()
    for epoch in range(dlstudio.epochs):

```

```

g_losses_per_print_cycle = []
d_losses_per_print_cycle = []
# For each batch in the dataloader
for i, data in enumerate(self.train_dataloader, 0):

    ## Maximization Part of the Min-Max Objective of Eq.
    (3):
        ##
        ## As indicated by Eq. (3) in the DCGAN part of the
        doc section at the beginning of this
        ## file, the GAN training boils down to carrying out
        a min-max optimization. Each iterative
        ## step of the max part results in updating the
        Discriminator parameters and each iterative
        ## step of the min part results in the updating of
        the Generator parameters. For each
        ## batch of the training data, we first do max and
        then do min. Since the max operation
        ## affects both terms of the criterion shown in the
        doc section, it has two parts: In the
        ## first part we apply the Discriminator to the
        training images using 1.0 as the target;
        ## and, in the second part, we supply to the
        Discriminator the output of the Generator
        ## and use 0 as the target. In what follows, the
        Discriminator is being applied to
        ## the training images:
        netD.zero_grad()
        real_images_in_batch =
data[0].to(self.dlstudio.device)
        # Need to know how many images we pulled in since at
        the tailend of the dataset, the
        # number of images may not equal the user-specified
        batch size:
        b_size = real_images_in_batch.size(0)
        label = torch.full((b_size,), real_label,
dtype=torch.float, device=self.dlstudio.device)
        output = netD(real_images_in_batch).view(-1)
        lossD_for_reals = criterion(output, label)

        lossD_for_reals.backward()

        ## That brings us the second part of what it takes to
        carry out the max operation on the
        ## min-max criterion shown in Eq. (3) in the doc
        section at the beginning of this file.
        ## part calls for applying the Discriminator to the
        images produced by the Generator from
        ## noise:

```



```

        noise = torch.randn(b_size, nz, 1, 1,
device=self.dlstudio.device)
        fakes = netG(noise)
        label.fill_(fake_label)
        ## The call to fakes.detach() in the next statement
returns a copy of the 'fakes' tensor
        ## that does not exist in the computational graph.
That is, the call shown below first
        ## makes a copy of the 'fakes' tensor and then
removes it from the computational graph.
        ## The original 'fakes' tensor continues to remain in
the computational graph. This ploy
        ## ensures that a subsequent call to backward() in
the 3rd statement below would only
        ## update the netD weights.
        output = netD(fakes.detach()).view(-1)
        lossD_for_fakes = criterion(output, label)
        lossD_for_fakes.backward()
        ## The following is just for displaying the losses:
        lossD = lossD_for_reals + lossD_for_fakes
        d_losses_per_print_cycle.append(lossD)
        ## Only the Discriminator weights are incremented:
        optimizerD.step()

        ## Minimization Part of the Min-Max Objective of Eq.
(3):
        ##
        ## That brings us to the min part of the max-min
optimization described in Eq. (3) the doc
        ## section at the beginning of this file. The min
part requires that we minimize
        ## "1 - D(G(z))" which, since D is constrained to lie
in the interval (0,1), requires that
        ## we maximize D(G(z)). We accomplish that by
applying the Discriminator to the output
        ## of the Generator and use 1 as the target for each
image:
        netG.zero_grad()
        label.fill_(real_label)
        ## The following forward prop will compute the
partials wrt the discriminator params also, but
        ## they will never get used for updating param vals
for two reasons: (1) We call "step()" on
        ## just optimizerG as shown later below; and (2) We
call "netD.zero_grad()" at the beginning of
        ## each training cycle.
        output = netD(fakes).view(-1)
        lossG = criterion(output, label)
        g_losses_per_print_cycle.append(lossG)

```

```

        lossG.backward()
        ## Only the Generator parameters are incremented:
        optimizerG.step()
        if i % 100 == 99:

            current_time = time.perf_counter()

            elapsed_time = current_time - start_time

            mean_D_loss =
torch.mean(torch.FloatTensor(d_losses_per_print_cycle))

            mean_G_loss =
torch.mean(torch.FloatTensor(g_losses_per_print_cycle))

            print("[epoch=%d/%d   iter=%4d   elapsed_time=%5d
secs]      mean_D_loss=%7.4f      mean_G_loss=%7.4f" %
                ((epoch+1), dlstudio.epochs,
(i+1), elapsed_time, mean_D_loss, mean_G_loss))
            d_losses_per_print_cycle = []

            g_losses_per_print_cycle = []

            G_losses.append(lossG.item())

            D_losses.append(lossD.item())

            if (iters % 500 == 0) or ((epoch == dlstudio.epochs-1)
and (i == len(self.train_dataloader)-1)):
                with torch.no_grad():
                    fake = netG(fixed_noise).detach().cpu() ##
detach() removes the fake from comp. graph. ##
for creating its CPU compatible version
                    img_list.append(torchvision.utils.make_grid(fake,
padding=1, pad_value=1, normalize=True))
                    iters += 1
            # At the end of training, make plots from the data in
G_losses and D_losses:
            plt.figure(figsize=(10,5))
            plt.title("Generator and Discriminator Loss During Training")

            plt.plot(G_losses, label="G")
            plt.plot(D_losses, label="D")
            plt.xlabel("iterations")
            plt.ylabel("Loss")
            plt.legend()
            plt.savefig(dir_name_for_results +
"/gen_and_disc_loss_training.png")
            plt.show()

```

```

        # Make an animated gif from the Generator output images
        stored in img_list:
        images = []
        for imgobj in img_list:
            img = tvtf.to_pil_image(imgobj)
            images.append(img)
        imageio.mimsave(dir_name_for_results +
"/generation_animation.gif", images, fps=5)
        # Make a side-by-side comparison of a batch-size sampling of
        real images drawn from the
        # training data and what the Generator is capable of
        producing at the end of training:
        real_batch = next(iter(self.train_dataloader))
        real_batch = real_batch[0]
        plt.figure(figsize=(15,15))
        plt.subplot(1,2,1)
        plt.axis("off")
        plt.title("Real Images")

plt.imshow(np.transpose(torchvision.utils.make_grid(real_batch.to(self
.dlstudio.device),
padding=1, pad_value=1,
normalize=True).cpu(),(1,2,0)))
        plt.subplot(1,2,2)

        plt.axis("off")

        plt.title("Fake Images")

        plt.imshow(np.transpose(img_list[-1],(1,2,0)))

        plt.savefig(dir_name_for_results + "/real_vs_fake_images.png")

        plt.show()

def run_wgan_code(self, dlstudio, critic, generator, results_dir):
    """
    This function is meant for training a CG1-based Critic-
    Generator WGAN. The implementation
    shown uses several programming constructs from the WGAN
    implementation at GitHub by the
    original authors of the famous WGAN paper. I have also used
    several programming constructs
    from the DCGAN code at PyTorch and GitHub. Regarding how to
    set the parameters of this method,
    see the following script in the "ExamplesAdversarialLearning"
    directory of the distribution:

        wgan_CG1.py

```

```

"""
dir_name_for_results = results_dir
if os.path.exists(dir_name_for_results):
    files = glob.glob(dir_name_for_results + "/*")
    for file in files:
        if os.path.isfile(file):
            os.remove(file)
        else:
            files = glob.glob(file + "/*")
            list(map(lambda x: os.remove(x), files))
else:
    os.mkdir(dir_name_for_results)
# Set the number of channels for the 1x1 input noise vectors
for the Generator:
nz = 100
netC = critic.to(self.dlstudio.device)
netG = generator.to(self.dlstudio.device)
# Initialize the parameters of the Critic and the Generator
networks according to the
# definition of the "weights_init()" method:
netC.apply(self.weights_init)
netG.apply(self.weights_init)
# We will use a the same noise batch to periodically check on
the progress made for the Generator:
fixed_noise = torch.randn(self.dlstudio.batch_size, nz, 1, 1,
device=self.dlstudio.device)
# These are for training the Critic, 'one' is for the part of
the training with actual
# training images, and 'minus_one' is for the part based on
the images produced by the
# Generator:
one = torch.FloatTensor([1]).to(self.dlstudio.device)
minus_one = torch.FloatTensor([-1]).to(self.dlstudio.device)
# Adam optimizers for the Critic and the Generator:
optimizerC = optim.Adam(netC.parameters(),
lr=dlstudio.learning_rate, betas=(self.beta1, 0.999))
optimizerG = optim.Adam(netG.parameters(),
lr=dlstudio.learning_rate, betas=(self.beta1, 0.999))
img_list = []
Gen_losses = []
Cri_losses = []
iters = 0
gen_iterations = 0
print("\n\nStarting Training Loop.....[Be very patient at
the beginning since the Critic must separately be taken through a few
hundred iterations of training before you get to see anything
displayed in your terminal window. Depending on your hardware, it may
take around 5 minutes. Subsequently, each 100 iterations will take
just a few seconds. ]\n\n")

```

```

start_time = time.perf_counter()
dataloader = self.train_dataloader
clipping_thresh = self.clipping_threshold
# For each epoch
for epoch in range(dlstudio.epochs):
    data_iter = iter(dataloader)
    i = 0
    ncritic = 5
    # As was stated in the WGAN part of the doc section for
the AdversarialLearning
    # class at the beginning of this file, a minimization of
the Wasserstein distance between
    # the distribution that describes the training data and
the distribution that has been learned
    # so far by the Generator can be translated into a
maximization of the difference of the
    # average outputs of a 1-Lipschitz function as applied to
the training images and as applied
    # to the output of the Generator. LEARNING THIS 1-
Lipschitz FUNCTION IS THE JOB OF THE CRITIC.
    # Since the Critic and the Generator parameters must be
updated independently, we start by
    # turning on the "requires_grad" property of the Critic
parameters:
    while i < len(dataloader):
        for p in netC.parameters():
            p.requires_grad = True
        if gen_iterations < 25 or gen_iterations % 500 == 0:
# the choices 25 and 500 are from WGAN
            ncritic = 100
            ic = 0
            ## The inner 'while' loop shown below calculates the
expectations in Eq. (8) in the doc section
            ## at the beginning of this file:
            while ic < ncritic and i < len(dataloader):
                ic += 1
                for p in netC.parameters():
                    p.data.clamp_(-clipping_thresh,
clipping_thresh)
                ## Training the Critic (Part 1):
                # The maximization needed for training the
Critic, as shown in Eq. (8) in the doc section
                # at the beginning of this file, consists of two
parts. The first part involves applying the
                # Critic network to just the training images,
with each image subject to a "gradient
                # target" of "-1".
                netC.zero_grad()

```

```

#                 real_images_in_batch = data_iter.next()
real_images_in_batch = next(data_iter)
i += 1
real_images_in_batch =
real_images_in_batch[0].to(self.dlstudio.device)
# Need to know how many images we pulled in since
at the tailend of the dataset, the
# number of images may not equal the user-
specified batch size:
b_size = real_images_in_batch.size(0)
# Note that a single scalar is produced for all
the data in a batch. This is probably
# the reason why what the Generator learns is
somewhat fuzzy.

critic_for_reals_mean = netC(real_images_in_batch)
## 'minus_one' is the gradient target:
critic_for_reals_mean.backward(minus_one)

## Training the Critic (Part 2):
# The second part of Critic training requires
that we apply the Critic to the images
# produced by the Generator for a fresh batch of
input noise vectors. The output of
# the Critic for these images must be subject to
the target "-1".

noise = torch.randn(b_size, nz, 1, 1,
device=self.dlstudio.device)
fakes = netG(noise)
# Again, a single number is produced for the
whole batch:

critic_for_fakes_mean = netC(fakes)
## 'one' is the gradient target:
critic_for_fakes_mean.backward(one)
wasser_dist = critic_for_reals_mean -
critic_for_fakes_mean
loss_critic = critic_for_fakes_mean -
critic_for_reals_mean
# Update the Critic
optimizerC.step()

## Training the Generator:
## That brings us to the training of the Generator
through the minimization required by the
## minmax objective in Eq. (7) at the beginning of
this file. To that end, first we must
## turn off the "requires_grad" of the Critic
parameters since the Critic and the Generator
## must be updated independently:
for p in netC.parameters():

```

```

        p.requires_grad = False
        netG.zero_grad()
        # This is again a single scalar based
        # characterization of the whole batch of the Generator images:
        noise = torch.randn(b_size, nz, 1, 1,
device=self.dlstudio.device)
        fakes = netG(noise)
        critic_for_fakes_mean = netC(fakes)
        loss_gen = critic_for_fakes_mean
        critic_for_fakes_mean.backward(minus_one)

        # Update the Generator
        optimizerG.step()

        gen_iterations += 1

        if i % (ncritic * 20) == 0:
            current_time = time.perf_counter()

            elapsed_time = current_time - start_time

            print("[epoch=%d/%d   i=%4d   el_time=%5d secs]
loss_critic=%7.4f   loss_gen=%7.4f   Wasserstein_dist=%7.4f" %
(epoch,dlstudio.epochs,i,elapsed_time,loss_critic.data[0],
loss_gen.data[0], wasser_dist.data[0]))
            Gen_losses.append(loss_gen.data[0].item())
            Cri_losses.append(loss_critic.data[0].item())
            # Get G's output on fixed_noise for the GIF
animation:
            if (iters % 500 == 0) or ((epoch == dlstudio.epochs-1)
and (i == len(dataloader)-1)):
                with torch.no_grad():

                    fake = netG(fixed_noise).detach().cpu() ##
detach() removes the fake from comp. graph.
                    ##
                    for its CPU compatible version
                        img_list.append(torchvision.utils.make_grid(fake,
padding=1, pad_value=1, normalize=True))
                        iters += 1

        # At the end of training, make plots from the data in
        Gen_losses and Cri_losses:
        plt.figure(figsize=(10,5))

        plt.title("Generator and Critic Loss During Training")

        plt.plot(Gen_losses,label="G")

```

```

plt.plot(Cri_losses, label="C")
plt.xlabel("iterations")
plt.ylabel("Loss")
plt.legend()

plt.savefig(dir_name_for_results +
"/gen_and_critic_loss_training.png")
plt.show()

# Make an animated gif from the Generator output images
stored in img_list:
images = []

for imgobj in img_list:
    img = tvtf.to_pil_image(imgobj)
    images.append(img)

imageio.mimsave(dir_name_for_results +
"/generation_animation.gif", images, fps=5)

# Make a side-by-side comparison of a batch-size sampling of
real images drawn from the
# training data and what the Generator is capable of
producing at the end of training:
real_batch = next(iter(dataloader))

real_batch = real_batch[0]
plt.figure(figsize=(15,15))

plt.subplot(1,2,1)

plt.axis("off")

plt.title("Real Images")

plt.imshow(np.transpose(torchvision.utils.make_grid(real_batch.to(self
.dlstudio.device),
padding=1, pad_value=1,
normalize=True).cpu(),(1,2,0)))
plt.subplot(1,2,2)

plt.axis("off")

plt.title("Fake Images")

```



```

plt.imshow(np.transpose(img_list[-1],(1,2,0)))

plt.savefig(dir_name_for_results + "/real_vs_fake_images.png")

plt.show()

def run_wgan_with_gp_code(self, dlstudio, critic, generator,
results_dir):
    """
        This function is meant for training a CG2-based Critic-
        Generator WGAN. Regarding how
        to set the parameters of this method, see the following script
        in the
        "ExamplesAdversarialLearning" directory of the distribution:

        wgan_with_gp_CG2.py
    """
    dir_name_for_results = results_dir
    if os.path.exists(dir_name_for_results):
        files = glob.glob(dir_name_for_results + "/*")
        for file in files:
            if os.path.isfile(file):
                os.remove(file)
            else:
                files = glob.glob(file + "/*")
                list(map(lambda x: os.remove(x), files))
    else:
        os.mkdir(dir_name_for_results)
        # Set the number of channels for the 1x1 input noise vectors
        for the Generator:
            nz = 100
            netC = critic.to(self.dlstudio.device)
            netG = generator.to(self.dlstudio.device)
            # Initialize the parameters of the Critic and the Generator
            networks according to the
            # definition of the "weights_init()" method:
            netC.apply(self.weights_init)
            netG.apply(self.weights_init)
            # We will use a the same noise batch to periodically check on
            the progress made for the Generator:
            fixed_noise = torch.randn(self.dlstudio.batch_size, nz, 1, 1,
            device=self.dlstudio.device)
            # These are for training the Critic, 'one' is for the part of
            the training with actual
            # training images, and 'minus_one' is for the part based on
            the images produced by the
            # Generator:
            one = torch.FloatTensor([1]).to(self.dlstudio.device)

```

```

        minus_one = torch.FloatTensor([-1]).to(self.dlstudio.device)
        # Adam optimizers for the Critic and the Generator:
        optimizerC = optim.Adam(netC.parameters(),
        lr=dlstudio.learning_rate, betas=(self.betal, 0.999))
        optimizerG = optim.Adam(netG.parameters(),
        lr=dlstudio.learning_rate, betas=(self.betal, 0.999))
        img_list = []
        Gen_losses = []
        Cri_losses = []
        iters = 0
        gen_iterations = 0
        start_time = time.perf_counter()
        dataloader = self.train_dataloader
        # For each epoch
        for epoch in range(dlstudio.epochs):
            data_iter = iter(dataloader)
            i = 0
            ncritic = 5
            # In this version of WGAN training, we enforce the 1-
            Lipschitz condition on the function
            # being learned by the Critic by requiring that the
            partial derivatives of the output of
            # the Critic with respect to its input equal one in
            magnitude. This is referred as imposing
            # a Gradient Penalty on the learning by the Critic. As
            in the previous training
            # function, we start by turning on the "requires_grad"
            property of the Critic parameters:
            while i < len(dataloader):
                for p in netC.parameters():
                    p.requires_grad = True
                ic = 0
                while ic < ncritic and i < len(dataloader):
                    ic += 1
                    # The first two parts of what it takes to train
                    the Critic are the same as for
                    # a regular WGAN. We want to train the Critic to
                    recognize the training images and,
                    # at the same time, the Critic should try to not
                    believe the output of the Generator.
                    netC.zero_grad()

            #
            real_images_in_batch = data_iter.next()
            real_images_in_batch = next(data_iter)
            i += 1
            real_images_in_batch =
            real_images_in_batch[0].to(self.dlstudio.device)
            # Need to know how many images we pulled in since
            at the tailend of the dataset, the

```

```

        # number of images may not equal the user-
specified batch size:
        b_size = real_images_in_batch.size(0)
        # Note that a single scalar is produced for all
the data in a batch.
        critic_for_reals_mean = netC(real_images_in_batch)
## this is a batch based mean
        # The gradient target is 'minus_one'. Note that
the gradient here is one of output of
        # the network with respect to the learnable
parameters:
        critic_for_reals_mean.backward(minus_one)
        # The second part of Critic training requires
that we apply the Critic to the images
        # produced by the Generator from a fresh batch of
input noise vectors.
        noise = torch.randn(b_size, nz, 1, 1,
device=self.dlstudio.device)
        fakes = netG(noise)
        # Again, a single number is produced for the
whole batch:
        critic_for_fakes_mean = netC(fakes.detach()) ##
detach() returns a copy of the 'fakes' tensor that has
##
been removed from the computational graph. This ensures
##
that a subsequent call to backward() will only update the Critic
        # The gradient target is 'one'. Note that the
gradient here is one of output of
        # the network with respect to the learnable
parameters:
        critic_for_fakes_mean.backward(one)
        # For the third part of Critic training, we need
to first estimate the Gradient Penalty
        # of the function being learned by the Critics
with respect to the input to the function.
        gradient_penalty =
self.calc_gradient_penalty(netC, real_images_in_batch, fakes)
        gradient_penalty.backward()
        loss_critic = critic_for_fakes_mean -
critic_for_reals_mean + gradient_penalty
        wasser_dist = critic_for_reals_mean -
critic_for_fakes_mean
        # Update the Critic
optimizerC.step()

        # That brings us to the training of the Generator.
First we must turn off the "requires_grad"
        # of the Critic parameters since the Critic and the

```

```

Generator are to be updated independently:
    for p in netC.parameters():
        p.requires_grad = False
    netG.zero_grad()
    # This is again a single scalar based
    characterization of the whole batch of the Generator images:
    noise = torch.randn(b_size, nz, 1, 1,
device=self.dlstudio.device)
    fakes = netG(noise)
    critic_for_fakes_mean = netC(fakes)
    loss_gen = critic_for_fakes_mean
    # The gradient target is 'minus_one'. Note that the
    gradient here is one of output of the network
    # with respect to the learnable parameters:
    loss_gen.backward(minus_one)
    # Update the Generator
    optimizerG.step()

    gen_iterations += 1
    if i % (ncritic * 20) == 0:
        current_time = time.perf_counter()

        elapsed_time = current_time - start_time

        print("[epoch=%d/%d   i=%4d   el_time=%5d secs]
loss_critic=%7.4f   loss_gen=%7.4f   Wasserstein_dist=%7.4f" %
(epoch+1,dlstudio.epochs,i,elapsed_time,loss_critic.data[0],
loss_gen.data[0], wasser_dist.data[0]))
        Gen_losses.append(loss_gen.data[0].item())
        Cri_losses.append(loss_critic.data[0].item())
        # Get G's output on fixed_noise for the GIF
    animation:
        if (iters % 500 == 0) or ((epoch == dlstudio.epochs-1)
and (i == len(self.train_dataloader)-1)):
            with torch.no_grad():

                fake = netG(fixed_noise).detach().cpu() ##
detach() removes the fake from comp. graph                                ## in
order to produce a CPU compatible tensor
                img_list.append(torchvision.utils.make_grid(fake,
padding=1, pad_value=1, normalize=True))
                iters += 1

        # At the end of training, make plots from the data in
        Gen_losses and Cri_losses:
        plt.figure(figsize=(10,5))

        plt.title("Generator and Critic Loss During Training")

```

```

plt.plot(Gen_losses, label="G")
plt.plot(Cri_losses, label="C")
plt.xlabel("iterations")
plt.ylabel("Loss")
plt.legend()

plt.savefig(dir_name_for_results +
"/gen_and_critic_loss_training.png")
plt.show()

# Make an animated gif from the Generator output images
stored in img_list:
images = []

for imgobj in img_list:
    img = tvtf.to_pil_image(imgobj)
    images.append(img)

imageio.mimsave(dir_name_for_results +
"/generation_animation.gif", images, fps=5)

# Make a side-by-side comparison of a batch-size sampling of
real images drawn from the
# training data and what the Generator is capable of
producing at the end of training:
real_batch = next(iter(self.train_dataloader))

real_batch = real_batch[0]
plt.figure(figsize=(15,15))

plt.subplot(1,2,1)

plt.axis("off")

plt.title("Real Images")

plt.imshow(np.transpose(torchvision.utils.make_grid(real_batch.to(self
.dlstudio.device),
padding=1, pad_value=1,
normalize=True).cpu(),(1,2,0)))
plt.subplot(1,2,2)

plt.axis("off")

```

```

plt.title("Fake Images")

plt.imshow(np.transpose(img_list[-1],(1,2,0)))

plt.savefig(dir_name_for_results + "/real_vs_fake_images.png")

plt.show()

def save_model(self, model):
    """
    Save the trained model to a disk file
    """
    torch.save(model.state_dict(),
self.dl_studio.path_saved_model)

```

1.2 Training curves

the code below is mainly borrowed from DLSudio-2.5.2

```

import random
import numpy
import torch
import cv2

seed = 0
random.seed(seed)
torch.manual_seed(seed)
torch.cuda.manual_seed(seed)
numpy.random.seed(seed)
torch.backends.cudnn.deterministic=True
torch.backends.cudnn.benchmark=False
os.environ['PYTHONHASHSEED'] = str(seed)

dls = DLStudio(
    dataroot = "./Supplementary/celeba_dataset_64x64/",
    # dataroot = "./dataGAN/PurdueShapes5GAN/multiobj/",

    image_size = [64,64],
    path_saved_model = "./saved_model",

    learning_rate = 2e-4,      ## <== try smaller value
    if mode collapse
    # learning_rate = 1e-5,      ## <== the output is

```

```

super bad
        epochs = 30,
        batch_size = 32,

        use_gpu = True,

    )

dcgan = MyAdversarialLearning(
    dlstudio = dls,
    ngpu = 1,
    latent_vector_size = 100,
    betal = 0.5,
Adam optimizer
)

discriminator = dcgan.DiscriminatorDG2()
generator = dcgan.GeneratorDG2()

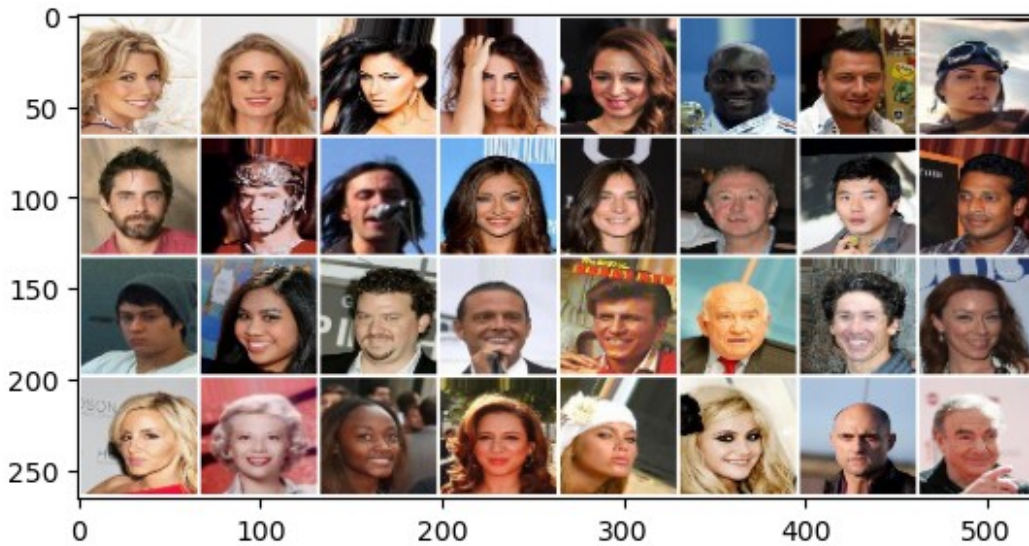
dcgan.set_data_loader()

dcgan.show_sample_images_from_dataset(dls)

dcgan.run_gan_code(dls, discriminator=discriminator,
generator=generator, results_dir="results_DG2")

# save model
torch.save(generator.state_dict(), "./generator.pth")
torch.save(discriminator.state_dict(), "./discriminator.pth")

```






Starting Training Loop...

```
[epoch=1/30  iter= 100  elapsed_time=  27 secs]    mean_D_loss=
0.2562    mean_G_loss=21.0271
[epoch=1/30  iter= 300  elapsed_time=  67 secs]    mean_D_loss=
0.7281    mean_G_loss= 9.4440
[epoch=2/30  iter= 100  elapsed_time=  99 secs]    mean_D_loss=
0.7063    mean_G_loss= 4.5943
[epoch=2/30  iter= 300  elapsed_time= 139 secs]    mean_D_loss=
0.7870    mean_G_loss= 4.5438
[epoch=3/30  iter= 100  elapsed_time= 172 secs]    mean_D_loss=
0.6686    mean_G_loss= 4.4085
[epoch=3/30  iter= 300  elapsed_time= 212 secs]    mean_D_loss=
0.7281    mean_G_loss= 4.5149
[epoch=4/30  iter= 100  elapsed_time= 245 secs]    mean_D_loss=
0.6891    mean_G_loss= 4.7506
[epoch=4/30  iter= 300  elapsed_time= 285 secs]    mean_D_loss=
0.6342    mean_G_loss= 4.7358
[epoch=5/30  iter= 100  elapsed_time= 319 secs]    mean_D_loss=
0.5143    mean_G_loss= 4.9340
[epoch=5/30  iter= 300  elapsed_time= 358 secs]    mean_D_loss=
0.4951    mean_G_loss= 5.1045
[epoch=6/30  iter= 100  elapsed_time= 392 secs]    mean_D_loss=
0.4403    mean_G_loss= 4.9777
[epoch=6/30  iter= 300  elapsed_time= 434 secs]    mean_D_loss=
0.5422    mean_G_loss= 5.0491
[epoch=7/30  iter= 100  elapsed_time= 490 secs]    mean_D_loss=
0.4655    mean_G_loss= 5.2195
[epoch=7/30  iter= 300  elapsed_time= 728 secs]    mean_D_loss=
```


0.5138	mean_G_loss= 4.8875		
[epoch=8/30	iter= 100	elapsed_time= 871 secs]	mean_D_loss=
0.4783	mean_G_loss= 5.1607		
[epoch=8/30	iter= 300	elapsed_time= 1110 secs]	mean_D_loss=
0.4826	mean_G_loss= 4.8877		
[epoch=9/30	iter= 100	elapsed_time= 1258 secs]	mean_D_loss=
0.4843	mean_G_loss= 4.6394		
[epoch=9/30	iter= 300	elapsed_time= 1497 secs]	mean_D_loss=
0.5271	mean_G_loss= 4.7227		
[epoch=10/30	iter= 100	elapsed_time= 1646 secs]	mean_D_loss=
0.6045	mean_G_loss= 4.7244		
[epoch=10/30	iter= 300	elapsed_time= 1885 secs]	mean_D_loss=
0.5134	mean_G_loss= 4.3287		
[epoch=11/30	iter= 100	elapsed_time= 2033 secs]	mean_D_loss=
0.5230	mean_G_loss= 4.2763		
[epoch=11/30	iter= 300	elapsed_time= 2272 secs]	mean_D_loss=
0.5189	mean_G_loss= 4.1503		
[epoch=12/30	iter= 100	elapsed_time= 2421 secs]	mean_D_loss=
0.5641	mean_G_loss= 4.2605		
[epoch=12/30	iter= 300	elapsed_time= 2659 secs]	mean_D_loss=
0.5540	mean_G_loss= 4.0672		
[epoch=13/30	iter= 100	elapsed_time= 2809 secs]	mean_D_loss=
0.4877	mean_G_loss= 4.1415		
[epoch=13/30	iter= 300	elapsed_time= 3048 secs]	mean_D_loss=
0.5645	mean_G_loss= 4.0690		
[epoch=14/30	iter= 100	elapsed_time= 3198 secs]	mean_D_loss=
0.4757	mean_G_loss= 3.8550		
[epoch=14/30	iter= 300	elapsed_time= 3436 secs]	mean_D_loss=
0.5281	mean_G_loss= 4.0452		
[epoch=15/30	iter= 100	elapsed_time= 3586 secs]	mean_D_loss=
0.4668	mean_G_loss= 4.0376		
[epoch=15/30	iter= 300	elapsed_time= 3825 secs]	mean_D_loss=
0.5711	mean_G_loss= 4.0152		
[epoch=16/30	iter= 100	elapsed_time= 3974 secs]	mean_D_loss=
0.4843	mean_G_loss= 4.0258		
[epoch=16/30	iter= 300	elapsed_time= 4213 secs]	mean_D_loss=
0.5655	mean_G_loss= 4.0645		
[epoch=17/30	iter= 100	elapsed_time= 4363 secs]	mean_D_loss=
0.5066	mean_G_loss= 4.0284		
[epoch=17/30	iter= 300	elapsed_time= 4601 secs]	mean_D_loss=
0.5583	mean_G_loss= 4.0195		
[epoch=18/30	iter= 100	elapsed_time= 4751 secs]	mean_D_loss=
0.5250	mean_G_loss= 3.8803		
[epoch=18/30	iter= 300	elapsed_time= 4990 secs]	mean_D_loss=
0.4980	mean_G_loss= 3.9769		
[epoch=19/30	iter= 100	elapsed_time= 5140 secs]	mean_D_loss=
0.4769	mean_G_loss= 4.0000		
[epoch=19/30	iter= 300	elapsed_time= 5379 secs]	mean_D_loss=
0.5529	mean_G_loss= 4.0067		

[epoch=20/30	iter= 100	elapsed_time= 5529 secs]	mean_D_loss=
0.4645	mean_G_loss= 3.9217		
[epoch=20/30	iter= 300	elapsed_time= 5768 secs]	mean_D_loss=
0.4621	mean_G_loss= 3.9671		
[epoch=21/30	iter= 100	elapsed_time= 5918 secs]	mean_D_loss=
0.4589	mean_G_loss= 4.0622		
[epoch=21/30	iter= 300	elapsed_time= 6157 secs]	mean_D_loss=
0.5181	mean_G_loss= 4.2106		
[epoch=22/30	iter= 100	elapsed_time= 6306 secs]	mean_D_loss=
0.3586	mean_G_loss= 4.1458		
[epoch=22/30	iter= 300	elapsed_time= 6545 secs]	mean_D_loss=
0.4391	mean_G_loss= 4.3382		
[epoch=23/30	iter= 100	elapsed_time= 6695 secs]	mean_D_loss=
0.6453	mean_G_loss= 4.4343		
[epoch=23/30	iter= 300	elapsed_time= 6934 secs]	mean_D_loss=
0.4301	mean_G_loss= 4.1801		
[epoch=24/30	iter= 100	elapsed_time= 7084 secs]	mean_D_loss=
0.4513	mean_G_loss= 4.4192		
[epoch=24/30	iter= 300	elapsed_time= 7323 secs]	mean_D_loss=
0.4361	mean_G_loss= 4.3060		
[epoch=25/30	iter= 100	elapsed_time= 7474 secs]	mean_D_loss=
0.4157	mean_G_loss= 4.2367		
[epoch=25/30	iter= 300	elapsed_time= 7713 secs]	mean_D_loss=
0.4173	mean_G_loss= 4.4900		
[epoch=26/30	iter= 100	elapsed_time= 7863 secs]	mean_D_loss=
0.3061	mean_G_loss= 4.4125		
[epoch=26/30	iter= 300	elapsed_time= 8102 secs]	mean_D_loss=
0.4041	mean_G_loss= 4.6156		
[epoch=27/30	iter= 100	elapsed_time= 8252 secs]	mean_D_loss=
0.4008	mean_G_loss= 4.6062		
[epoch=27/30	iter= 300	elapsed_time= 8491 secs]	mean_D_loss=
0.3642	mean_G_loss= 4.5314		
[epoch=28/30	iter= 100	elapsed_time= 8641 secs]	mean_D_loss=
0.4097	mean_G_loss= 4.5341		
[epoch=28/30	iter= 300	elapsed_time= 8880 secs]	mean_D_loss=
0.3419	mean_G_loss= 4.5949		
[epoch=29/30	iter= 100	elapsed_time= 9029 secs]	mean_D_loss=
0.3106	mean_G_loss= 4.6599		
[epoch=29/30	iter= 300	elapsed_time= 9268 secs]	mean_D_loss=
0.4008	mean_G_loss= 4.7462		
[epoch=30/30	iter= 100	elapsed_time= 9418 secs]	mean_D_loss=
0.3484	mean_G_loss= 4.7665		
[epoch=30/30	iter= 300	elapsed_time= 9657 secs]	mean_D_loss=
0.1682	mean_G_loss= 4.6479		



```
# use genertor to generate images and save them for further use
gen = generator.to(dls.device)
gen.eval()
with torch.no_grad():

    # we need 1024 in total, so we need loop 32 times
    for j in range(32):
        noise = torch.randn(32, 100, 1, 1, device=dls.device)
        fake = gen(noise).detach().cpu() # shape: (32, C, H, W)

        # save them in output dir as png files
        os.makedirs("Generator_Output", exist_ok=True)

        for i in range(fake.size(0)): # 32 in this case
            img = fake[i] # shape: (C, H, W)
            img = img.numpy().transpose(1, 2, 0) # to (H, W, C)

            img = (img + 1) / 2.0 # Normalize from [-1,1] → [0,1]
            img = np.clip(img * 255.0, 0, 255) # Scale to [0,255]
```

```

img = img.astype(np.uint8)

# Convert from RGB to BGR (OpenCV uses BGR format)
img = cv2.cvtColor(img, cv2.COLOR_RGB2BGR)

cv2.imwrite(f"Generator_Output/fake_{j*32 + i}.jpg", img)

```

2 Diffusion

2.1 Generate Images

Since my personal laptop GPU isn't powerful enough to run multiple epochs of diffusion training with a batch size of 32, I am using the provided weights and skipping the execution of RunCodeForDiffusion.py.

```

# the code below is mainly borrowed from DLSudio-2.5.2
import numpy as np

GenerativeDiffusion_dir = os.path.join(DLStudio_dir,
"GenerativeDiffusion")
sys.path.append(GenerativeDiffusion_dir)

from GenerativeDiffusion import *

gauss_diffusion = GaussianDiffusion(
    num_diffusion_timesteps = 1000,
)

network = UNetModel(
    in_channels=3,
    model_channels = 128,
    out_channels = 3,
    num_res_blocks = 2,
    attention_resolutions = (4, 8),
## for 64x64 images
    channel_mult = (1, 2, 3, 4),
## for 64x64 images
    num_heads = 1,
    attention = True ## <<<
Must be the same as for RunCodeForDiffusion.py
    # attention = False ## <<<
Must be the same as for RunCodeForDiffusion.py
)

```

```

top_level = GenerativeDiffusion(
    gen_new_images      = True,
    image_size          = 64,
    num_channels         = 128,
    ema_rate            = 0.9999,
    diffusion = gauss_diffusion,
    network = network,
    ngpu = 1,
    path_saved_model = "./RESULTS",
    clip_denoised=True,
    num_samples=1024,
    batch_size_image_generation=8,
)

model_path = "Supplementary/diffusion.pt"

network.load_state_dict( torch.load(model_path) )

network.to(top_level.device)
network.eval()

print("sampling...")
all_images = []

while len(all_images) * top_level.batch_size_image_generation <
top_level.num_samples:
    sample = gauss_diffusion.p_sampler_for_image_generation(
        network,
        (top_level.batch_size_image_generation, 3,
top_level.image_size, top_level.image_size),
        device = top_level.device,
        clip_denoised = top_level.clip_denoised,
    )
    sample = ((sample + 1) * 127.5).clamp(0, 255).to(torch.uint8)
    sample = sample.permute(0, 2, 3, 1)
    sample = sample.contiguous()
    gathered_samples = [sample]
    all_images.extend([sample.cpu().numpy() for sample in
gathered_samples])
    print(f"created {len(all_images) *
top_level.batch_size_image_generation} samples")

arr = np.concatenate(all_images, axis=0)
arr = arr[: top_level.num_samples]

```

```

shape_str = "x".join([str(x) for x in arr.shape])
out_path = os.path.join("./", f"samples_{shape_str}.npz")
np.savez(out_path, arr)

print("image generation completed")

from PIL import Image
import os
import numpy as np
import glob
import matplotlib.pyplot as plt
import torch
import torchvision

npz_archive = "./samples_1024x64x64x3.npz"      ## You must change
this as needed. The first number, 32,         ## is the
batch-size used for sampling a checkpoint. It  ## is set in
the script GenerateNewImageSamples.py.         ##
an npz archive is a dict with a single         ## In our case,
pair. The name of the 'key' is 'arr_0'.         ## 'key-value'
shape of the 'value' will be the shape of the  ## And the
stores the generated images. For the example   ## ndarray that
shape of the value is (32,64,64,3)             ## shown, the
visualization_dir = "Diffusion_Output"
# image_display_size = (256,256)

if os.path.exists(visualization_dir):
    files = glob.glob(visualization_dir + "/*")
    for file in files:
        if os.path.isfile(file):
            os.remove(file)
else:
    os.mkdir(visualization_dir)

data = np.load(npz_archive)

for i, arr in enumerate(data['arr_0']):

```


```
img = Image.fromarray( arr )  
# img = img.resize( image_display_size )  
img.save( f"Diffusion_Output/test_{i}.jpg" )
```

3 FID score

3.1 Display 4x4(16) images generated by GAN

```
from PIL import Image  
def display_images(img_dir):  
    img_files = os.listdir(img_dir)  
    img_files = img_files[:16] # get the first 16 images  
  
    fig, axes = plt.subplots(4, 4, figsize=(10, 10))  
    fig.suptitle(img_dir)  
  
    for axis, img_file in zip(axes.flatten(), img_files): #  
axes.flatten() → [ax0, ax1, ..., ax15]. zip → [(ax0, img0), (ax1, img1), ...]  
        img_path = os.path.join(img_dir, img_file)  
        img = Image.open(img_path)  
        axis.imshow(img)  
        axis.axis('off')  
  
    plt.tight_layout()  
    plt.show()  
  
img_dir = "Generator_Output"  
display_images(img_dir)
```


Generator_Output



3.2 Display 4x4(16) images generated by Diffusion

```
img_dir = "Diffusion_Output\Diffusion_Output"  
display_images(img_dir)
```


Diffusion_Output\Diffusion_Output



3.3 GAN FID

```
from pytorch_fid.fid_score import calculate_activation_statistics,
calculate_frechet_distance
from pytorch_fid.inception import InceptionV3

device = 'cuda' if torch.cuda.is_available() else 'cpu'

def calculate_fid(real_img_dir, fake_img_dir):
    dims = 2048
```

```

block_idx = InceptionV3.BLOCK_INDEX_BY_DIM[dims]
model = InceptionV3([block_idx]).to(device)

# make images in the directory as numpy arrays
fake_img_files = os.listdir(fake_img_dir)
real_img_files = os.listdir(real_img_dir)

fake_paths = [os.path.join(fake_img_dir, img_file) for img_file in
fake_img_files]
real_paths = [os.path.join(real_img_dir, img_file) for img_file in
real_img_files]

m1, s1 = calculate_activation_statistics(real_paths, model,
device=device)
m2, s2 = calculate_activation_statistics(fake_paths, model,
device=device)

fid_value = calculate_frechet_distance(m1, s1, m2, s2)

return fid_value

fake_img_dir = "Generator_Output"
real_img_dir =
"Supplementary/celeba_dataset_64x64/celeba_dataset_64x64"

# # calculate FID score
fid_value = calculate_fid(real_img_dir, fake_img_dir)
print(f"FID score of GAN: {fid_value}")

Using device: cuda
FID score of GAN: 66.20835395075642

```

3.4 Diffusion FID

```

fake_img_dir = "Diffusion_Output"

fid_value = calculate_fid(real_img_dir, fake_img_dir)
print(f"FID score of Diffusion: {fid_value}")

FID score of Diffusion: 55.308944348883074

```

- Qualitative Evaluation:

For the GAN-generated (**Generator**) images, there is a noticeable lack of detail — textures appear vague and poorly defined. Additionally, the facial proportions are often inconsistent, with some features being unnaturally placed or distorted. However, in my generated dataset, I did not observe signs of mode collapse — the samples appear diverse and vary across different runs.

In contrast, the images generated by the Diffusion model (**p-chain**) are significantly sharper and more realistic. The textures are more consistent and natural-looking, and facial features are generally well-structured and proportionally accurate. Overall, the visual quality of the Diffusion outputs is noticeably higher than that of the GAN.

One reason Diffusion models tend to outperform GANs is due to the fundamental difference in how they are trained and how they generate data. GANs start from Gaussian noise and learn a direct mapping to output images in a single forward pass, which makes training unstable and prone to issues like mode collapse or artifacts. Diffusion models, on the other hand, follow a gradual denoising process: they start with pure noise and iteratively remove noise over many steps to generate a high-quality image. This slow, step-by-step generation allows diffusion models to capture fine-grained details and complex data distributions more effectively.

Architecturally, GANs rely heavily on the adversarial setup between a generator and a discriminator, which often creates a delicate balance that's hard to maintain. Diffusion models, in contrast, are trained with a denoising score matching objective and don't rely on adversarial loss — making training more stable and less sensitive to hyperparameters.

- Quantitative Evaluation:

	GAN	Diffusion
FID Score	66.2	55.3

Although visually Diffusion has better results, the Diffusion model scored 55.3 in FID, while the GAN scored 66.2. I think the reason for this discrepancy lies in how FID evaluates similarity to the real data distribution. Despite the higher visual quality of Diffusion samples, they may introduce subtle distributional shifts or over-smooth textures that deviate statistically from the real images, which can affect the FID score. On the other hand, GANs — while sometimes producing artifacts or less detailed images — might still align more closely with the overall statistical moments (mean and covariance) of the dataset, thus receiving a slightly better FID score. This highlights a known limitation of FID: it doesn't always perfectly correlate with human perception of image quality. Therefore, a combination of both qualitative and quantitative evaluation provides a more complete picture of generative model performance.

Finetune GAN

4.1 Training curves

the code below is mainly borrowed from DLSudio-2.5.2

```
import random
import numpy
import torch
import cv2

seed = 0
random.seed(seed)
torch.manual_seed(seed)
```

```

torch.cuda.manual_seed(seed)
numpy.random.seed(seed)
torch.backends.cudnn.deterministic=True
torch.backends.cudnn.benchmark=False
os.environ['PYTHONHASHSEED'] = str(seed)

dls = DLStudio(
    dataroot = "./Diffusion_Output/",
    image_size = [64,64],
    path_saved_model = "./saved_model",
    # learning_rate = 2e-4,      ## <== mode collapse
    # learning_rate = 2e-5,      ## <== mode collapse
    learning_rate = 1e-5,      ## <== mode collapse
    epochs = 30,
    batch_size = 32,
    use_gpu = True,
)

dcgan = MyAdversarialLearning(
    dlstudio = dls,
    ngpu = 1,
    latent_vector_size = 100,
    betal = 0.5,                ## for the
    Adam optimizer
)

discriminator = dcgan.DiscriminatorDG2()

# using the pretrained model make the discriminator too strong, it
# gives very harsh gradients to the generator.
# discriminator.load_state_dict(torch.load("./discriminator.pth"))
# discriminator.to(dls.device)

generator = dcgan.GeneratorDG2()
generator.load_state_dict(torch.load("./generator.pth"))
generator.to(dls.device)

dcgan.set_data_loader()

dcgan.show_sample_images_from_dataset(dls)

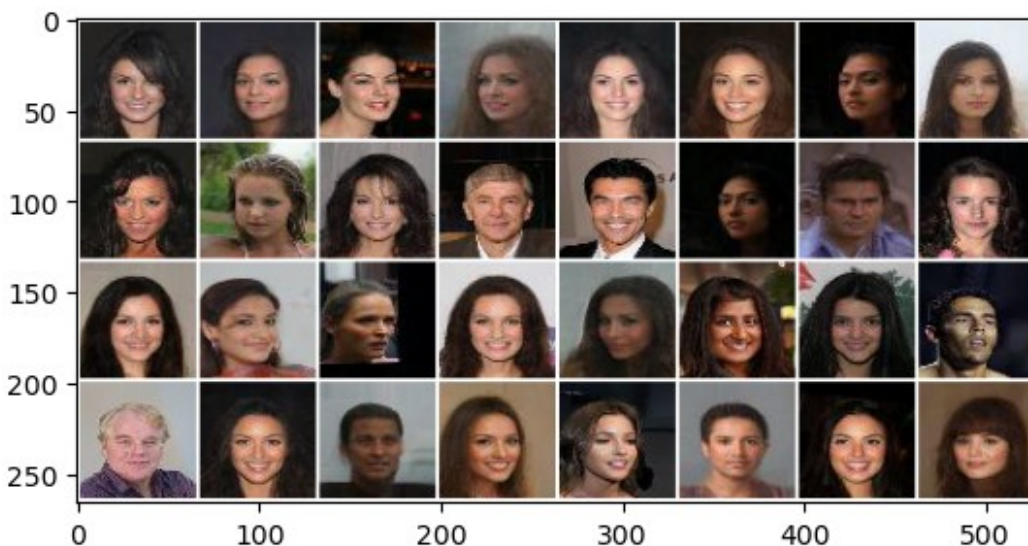
dcgan.run_gan_code(dls, discriminator=discriminator,
generator=generator, results_dir="results_DG2_finetuned")

```

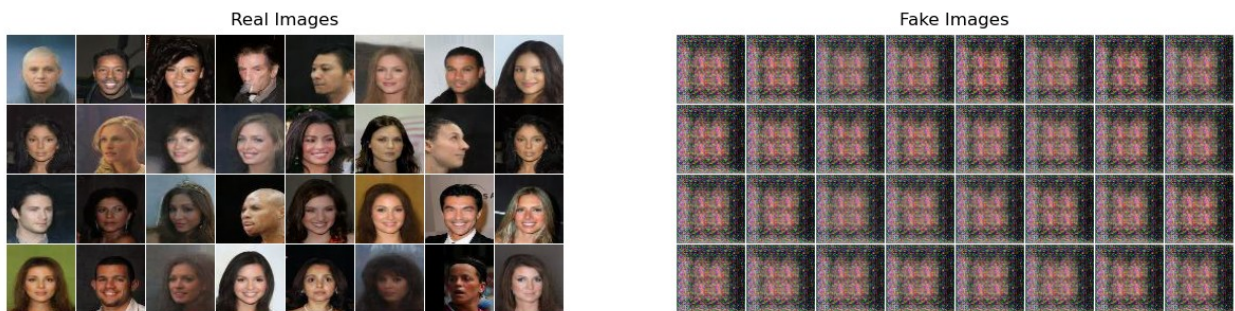

```
# save model
```

```
torch.save(generator.state_dict(), "./generator_finetuned.pth")  
torch.save(discriminator.state_dict(),  
"./discriminator_finetuned.pth")
```

```
C:\Users\alanc\AppData\Local\Temp\ipykernel_11212\377586184.py:42:  
FutureWarning: You are using `torch.load` with `weights_only=False`  
(the current default value), which uses the default pickle module  
implicitly. It is possible to construct malicious pickle data which  
will execute arbitrary code during unpickling (See  
https://github.com/pytorch/pytorch/blob/main/SECURITY.md#untrusted-models  
for more details). In a future release, the default value for  
`weights_only` will be flipped to `True`. This limits the functions  
that could be executed during unpickling. Arbitrary objects will no  
longer be allowed to be loaded via this mode unless they are  
explicitly allowlisted by the user via  
`torch.serialization.add_safe_globals`. We recommend you start setting  
`weights_only=True` for any use case where you don't have full control  
of the loaded file. Please open an issue on GitHub for any issues  
related to this experimental feature.  
generator.load_state_dict(torch.load("./generator.pth"))
```



Starting Training Loop...



4.2 Visual observations of finetuned GAN against GAN

The outputs of the fine-tuned GAN show classic mode collapse — every generated image is an identical, abstract pattern that does not resemble human faces. Despite experimenting with various learning rates and even removing the pretrained discriminator (allowing it to learn from scratch), mode collapse persisted across runs.

In fact, the performance of the fine-tuned GAN is worse than the original GAN. This degradation likely stems from the fact that the fine-tuned model was trained on diffusion-generated images, which, while high in visual quality, are still synthetic. These images can carry subtle biases, overly smooth textures, or a lack of true diversity compared to real-world data, leading to poor generalization when used to train a GAN.

4.3 Visual observations of finetuned GAN against Diffusion

The diffusion model continues to produce sharp, diverse, and realistic human faces, while the fine-tuned GAN fails to do so. This contrast emphasizes how much better the diffusion model captures data distributions, even compared to a GAN trained (or fine-tuned) on diffusion outputs.

4.4 Display 4x4(16) images generated by finetuned GAN

```
# use genertor to generate images and save them for further use
gen = generator.to(dls.device)
gen.eval()
with torch.no_grad():
    # we need 1024 in total, so we need loop 32 times
    for j in range(32):
        noise = torch.randn(32, 100, 1, 1, device=dls.device)
        fake = gen(noise).detach().cpu() # shape: (32, C, H, W)

        # save them in output dir as png files
        os.makedirs("Generator_finetuned_Output", exist_ok=True)

        for i in range(fake.size(0)): # 32 in this case
            img = fake[i] # shape: (C, H, W)
            img = img.numpy().transpose(1, 2, 0) # to (H, W, C)


            img = (img + 1) / 2.0 # Normalize from [-1,1] → [0,1]
            img = np.clip(img * 255.0, 0, 255) # Scale to [0,255]
            img = img.astype(np.uint8)

            # Convert from RGB to BGR (OpenCV uses BGR format)
            img = cv2.cvtColor(img, cv2.COLOR_RGB2BGR)

            cv2.imwrite(f"Generator_finetuned_Output/fake_{j*32 +
i}.jpg", img)

img_dir = "Generator_finetuned_Output"
display_images(img_dir)
```

GAN Generated Images



- Discuss how training with diffusion-generated images affects the GANs performance.

Training a GAN on diffusion-generated images instead of real data significantly hurts performance. Visually, the generator fails to capture human-like structure and instead learns to produce repetitive, abstract patterns — an indication that it's not extracting meaningful or generalizable features. The root issue is **distribution mismatch**.

Diffusion-generated images, though realistic-looking, are still synthetic and may contain distributional artifacts, lack the complex diversity found in natural data, and be too smooth or consistent, offering fewer cues for the discriminator

As a result, the discriminator struggles to learn meaningful boundaries, weakening the training signal. Additionally, the generator overfits to superficial patterns present in the diffusion images. Finally, the model enters mode collapse, repeatedly generating the same output.